EMR on AWS

Elastic MapReduce on Amazon Web Services

Cosmin Stamate

Hadoop Streaming

- Streaming application reads input from *standard input* and then runs a script or executable (called a mapper) against each input.
- The result from each of the inputs is saved locally, typically on a Hadoop Distributed File System (HDFS) partition.
- After the input is processed by the mapper, a second script, called a reducer processes the mapper results.
- The results from the reducer are sent to *standard output*. You can chain together a series of Streaming steps, where the output of one step becomes the input of another step.

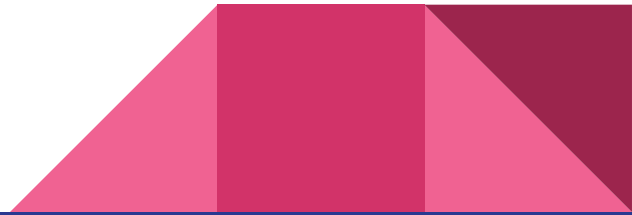

Hadoop Streaming

- Easier to implement, allows you to focus more on algorithm design and less on re-creating boilerplate code.
- Allows for easy local testing of MapReduce applications
- On AWS EMR we can write MapReduce applications in many languages if we use the streaming program interface.
- We can code mappers, reducers and combiners, not only Java, but also in other languages like Python, Perl, Ruby, PHP, or Bash.
- The only essential thing to remember is that we are using *standard input* and *standard output* to feed our MapReduce streaming functions.

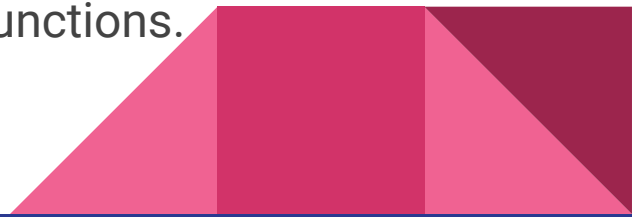

Datasets

Please download the following books in plain text format, which have been sourced from the Gutenberg Project

- http://www.dcs.bbk.ac.uk/~cosmin/cc/data/pg27827.txt
- http://www.dcs.bbk.ac.uk/~cosmin/cc/data/pg3207.txt
- http://www.dcs.bbk.ac.uk/~cosmin/cc/data/pg5200.txt

And the following google 1-grams which have been sourced from the Google Books Ngram Viewer

- http://www.dcs.bbk.ac.uk/~cosmin/cc/data/ngrams.txt

Mappers and reducers

We will be doing a word count on the books downloaded from the gutenberg project. We will use only a mapper and the **aggregate** function *(If you look in the wordcount-map.py file you will see that we are using the **LongValueSum** aggregate function)* instead of a actual reducer:

- http://www.dcs.bbk.ac.uk/~cosmin/cc/wordcount-map.py

We are going to use google ngrams (1-gram) to look for words which were coined in the year 1999

- http://www.dcs.bbk.ac.uk/~cosmin/cc/ngram-map.py
- http://www.dcs.bbk.ac.uk/~cosmin/cc/ngram-reduce.py

NOTE: python mapper and reducer

Please note that if you are using the Streaming API and are writing python mappers and reducers you have to add **#!/usr/bin/python** at the top of your python scripts as this is the path to the executable python. This needs to be there for all the scripts you have. The same applies for all other languages, for example if you will use ruby please add **#!/usr/bin/ruby** at the top of your script.

You can also use **#!/usr/bin/env python** depending on your environment setup.

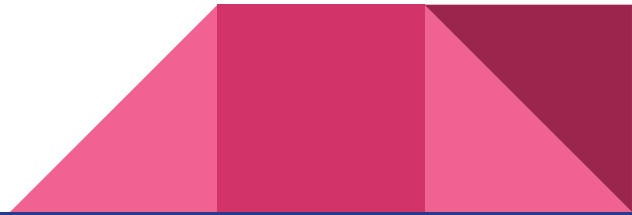

Create a free **AWS** account

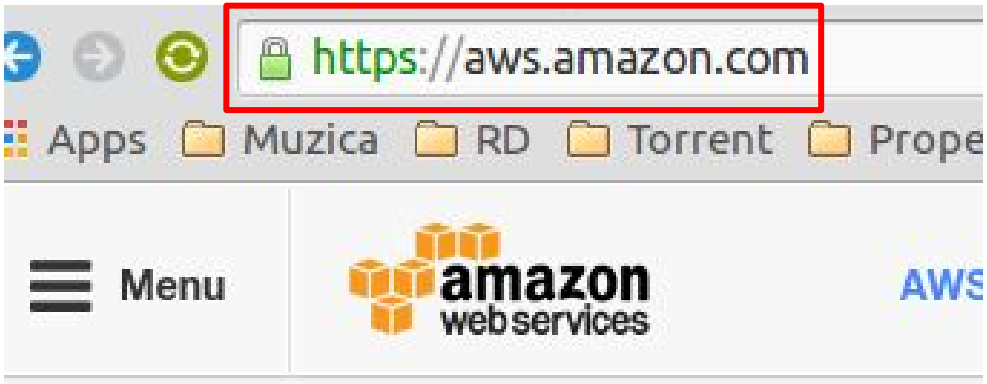

Apply for the Academic discount

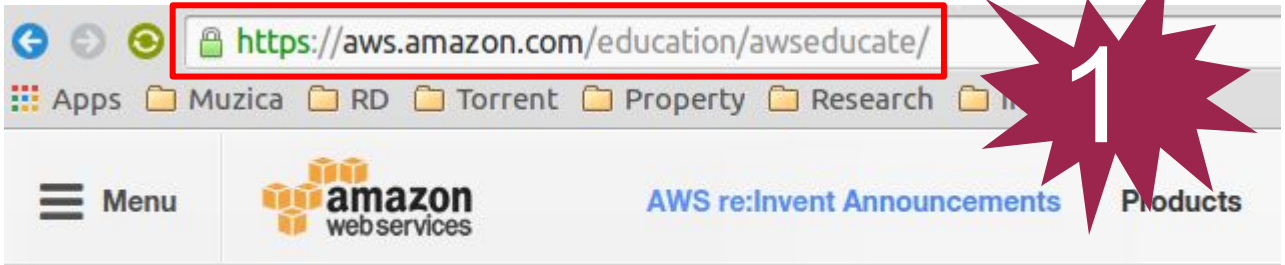

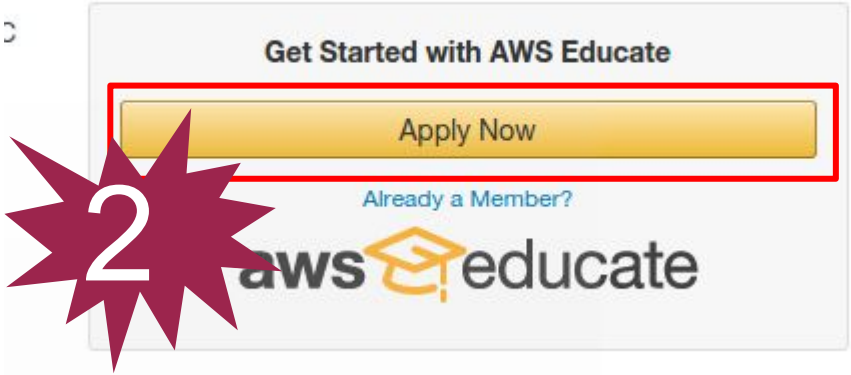

Fill in the form and use your university email

- You can get your **AWS account ID** (12 digit number) by loggin in to your AWS console and going to **My Account**, under your name.

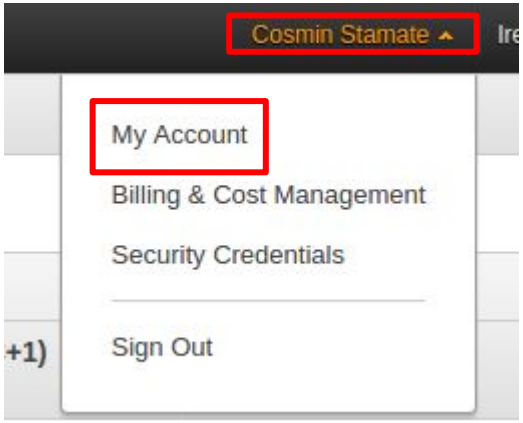

After you receive the **AWS Educate** Application Approved email

- ❏ Go to **My Account > Credits**
- ❏ Paste the promo-code from the approval email and redeem the credits

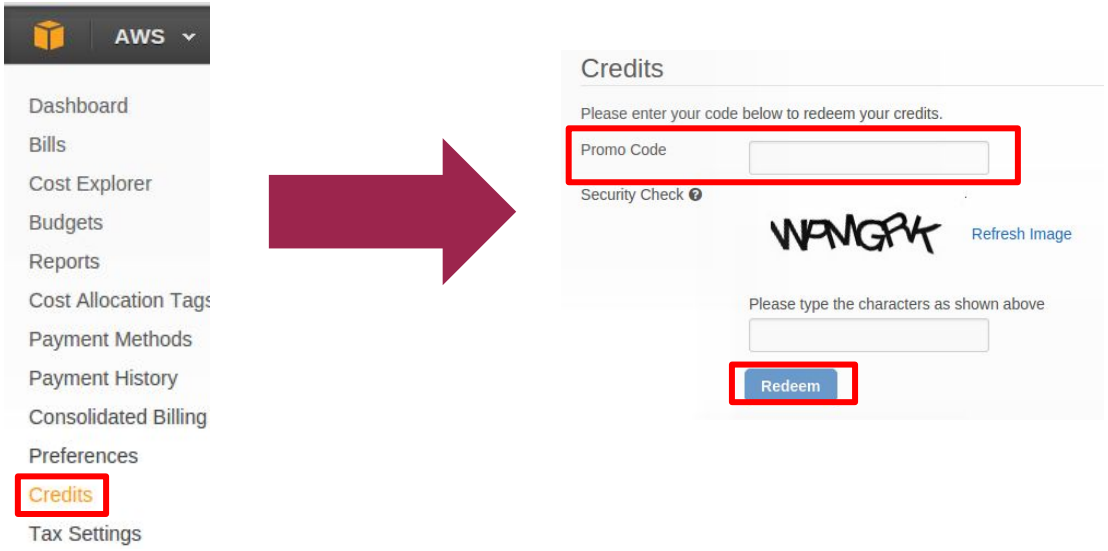

Congratulations, you now have $100 credits!

Expiration Date	Credit Name	Credits Used	Credits Remaining
2016-09-30	ENG_FY2015_Q4_100USD	$0.00	$100.00

Total Amount of Credits Remaining: $100.00

Sign in to the AWS Console

Chose **EU (Ireland)** region

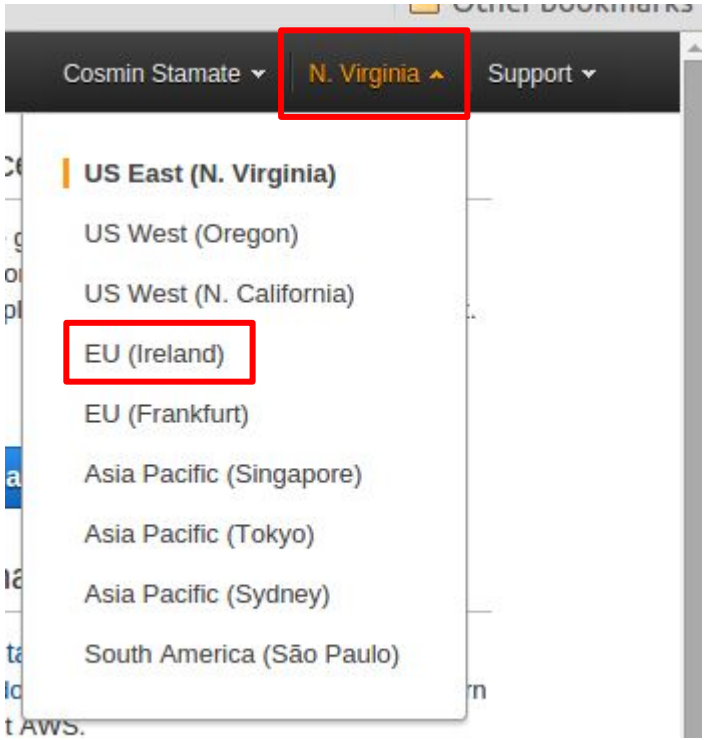

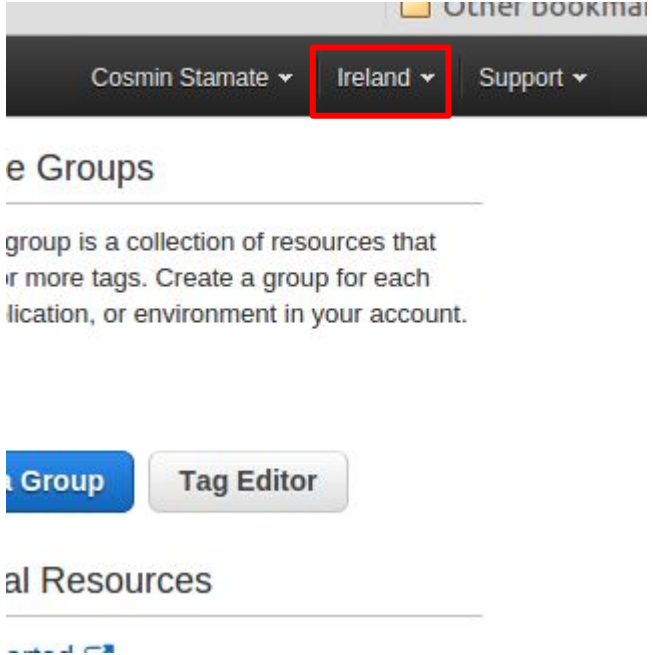

Create a **bucket** under **Storage > S3**

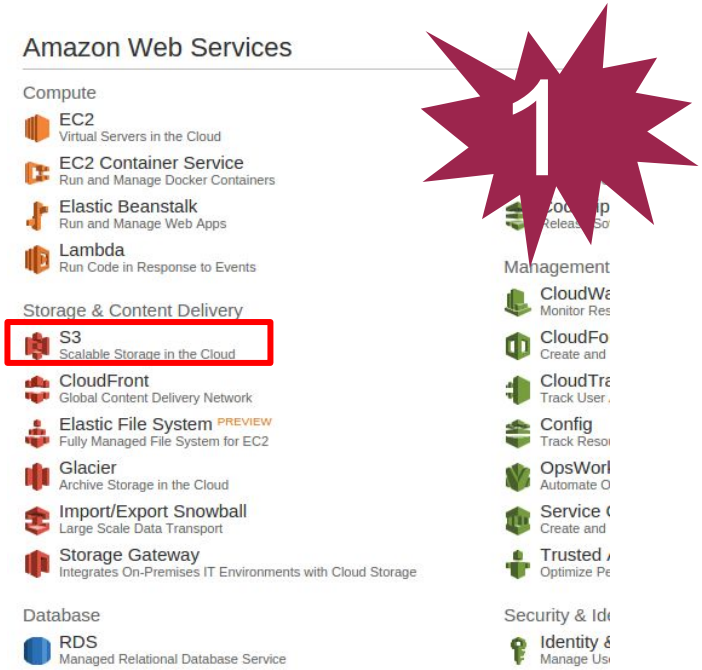

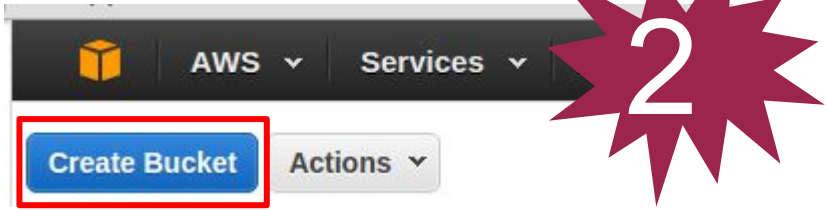

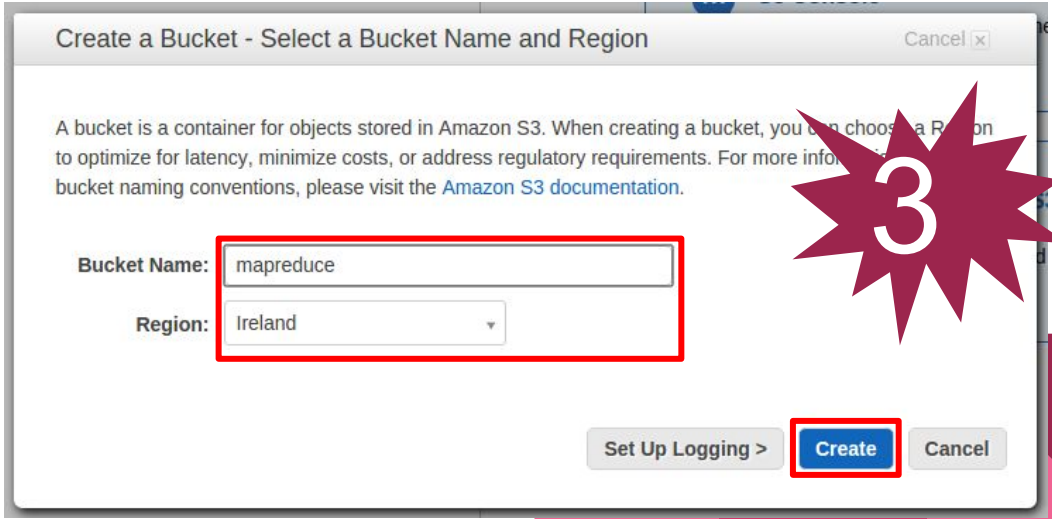

Add folders to the newly created **S3 bucket**

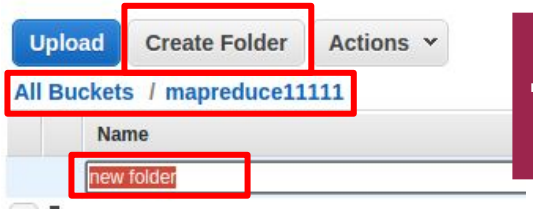

four folders

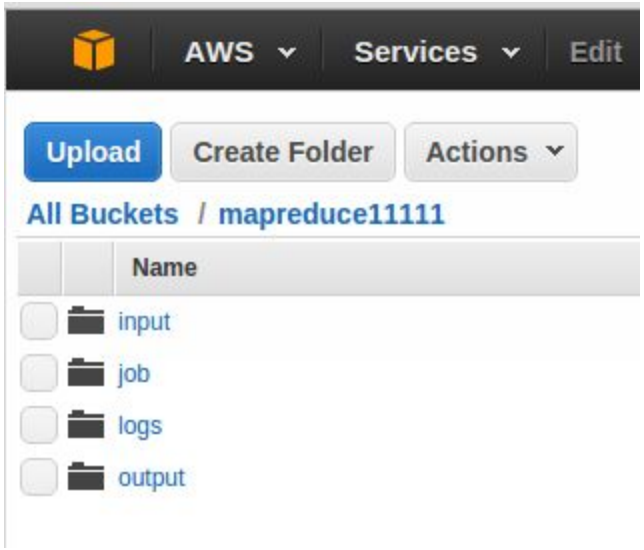

Upload wordcount books into their **S3 bucket** folder

Select the input folder and press Upload

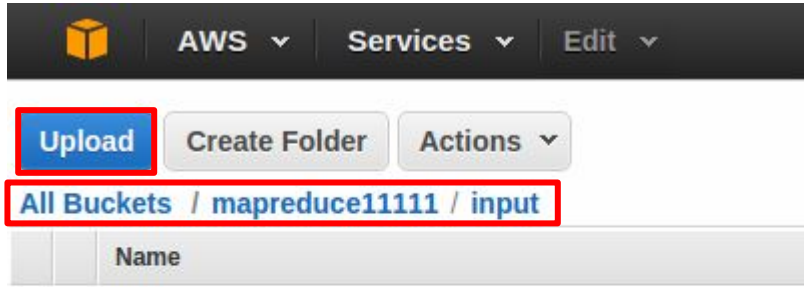

Add all your input files and upload

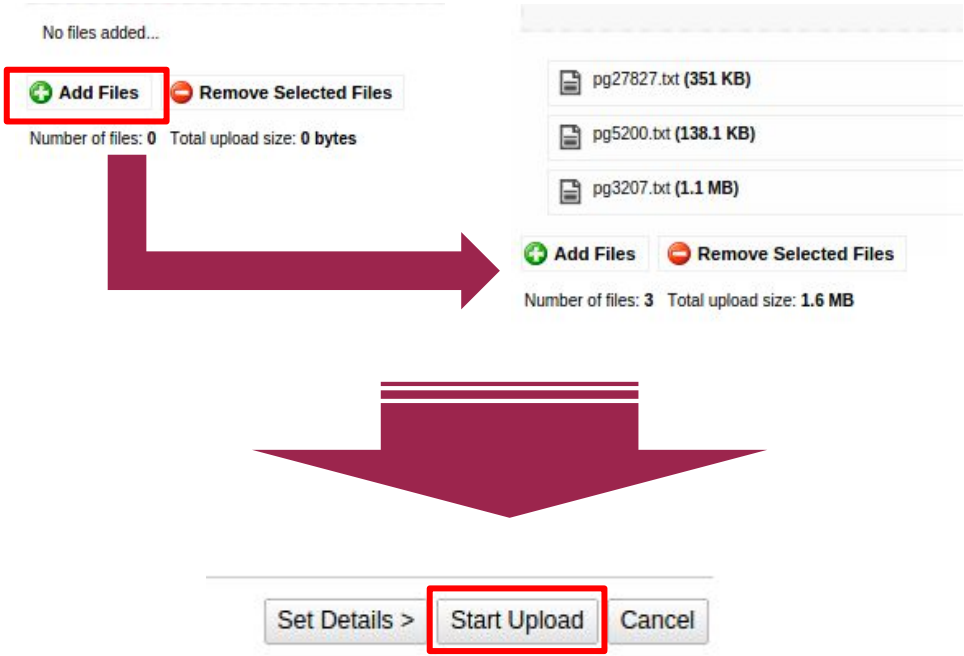

Upload your wordcount MapReduce files

Go to your **job** folder and press upload

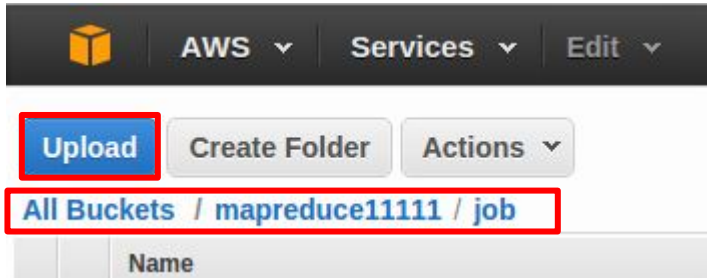

Select your files and upload them

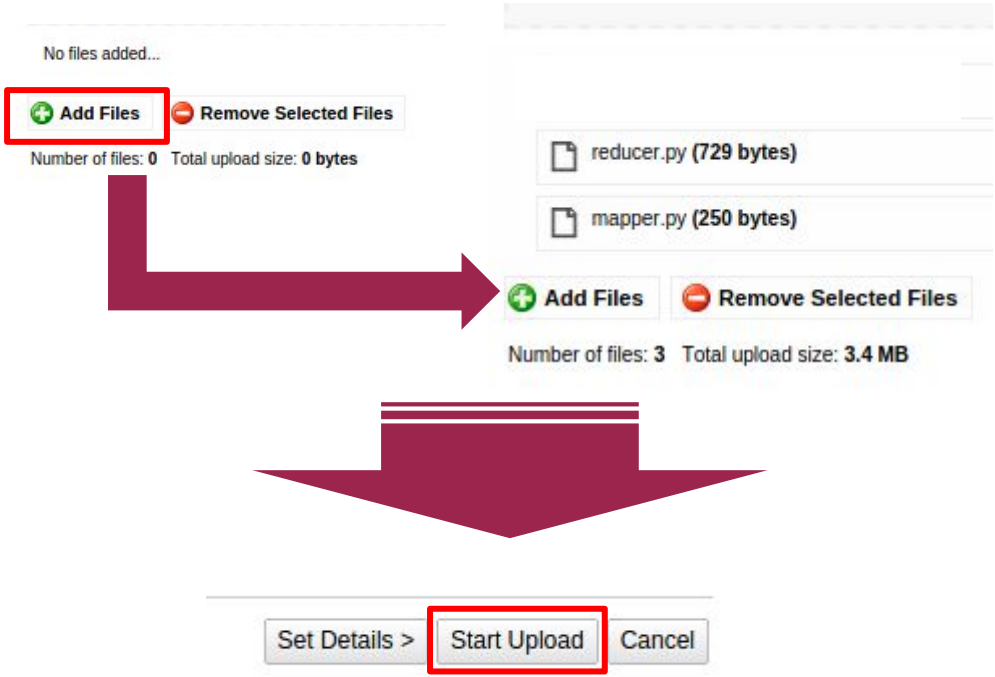

Create an **EMR cluster**

Under Analytics, press on **EMR**

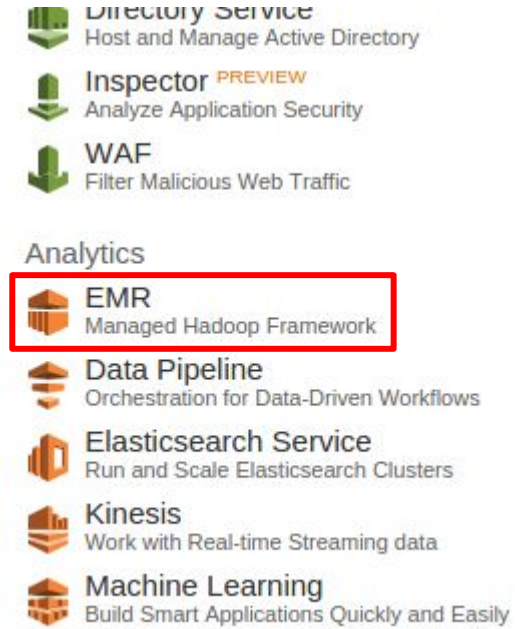

On the next page, press **Create cluster**

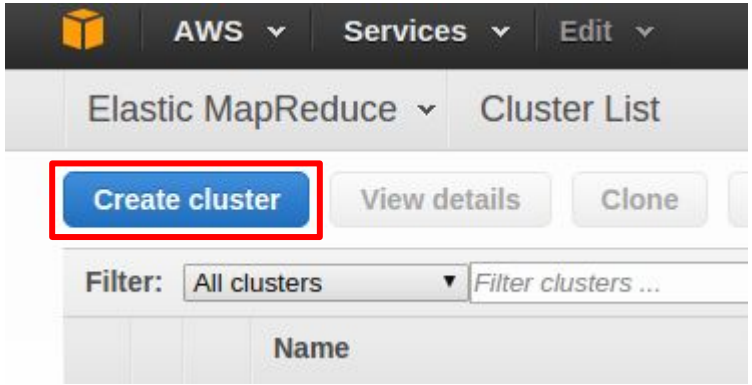

Create Cluster - Quick Options

- ❑ Choose a unique cluster name
- ❑ Enable logging and select the **logs/** S3 folder that you created earlier
- ❑ Choose step execution under **Launch mode**

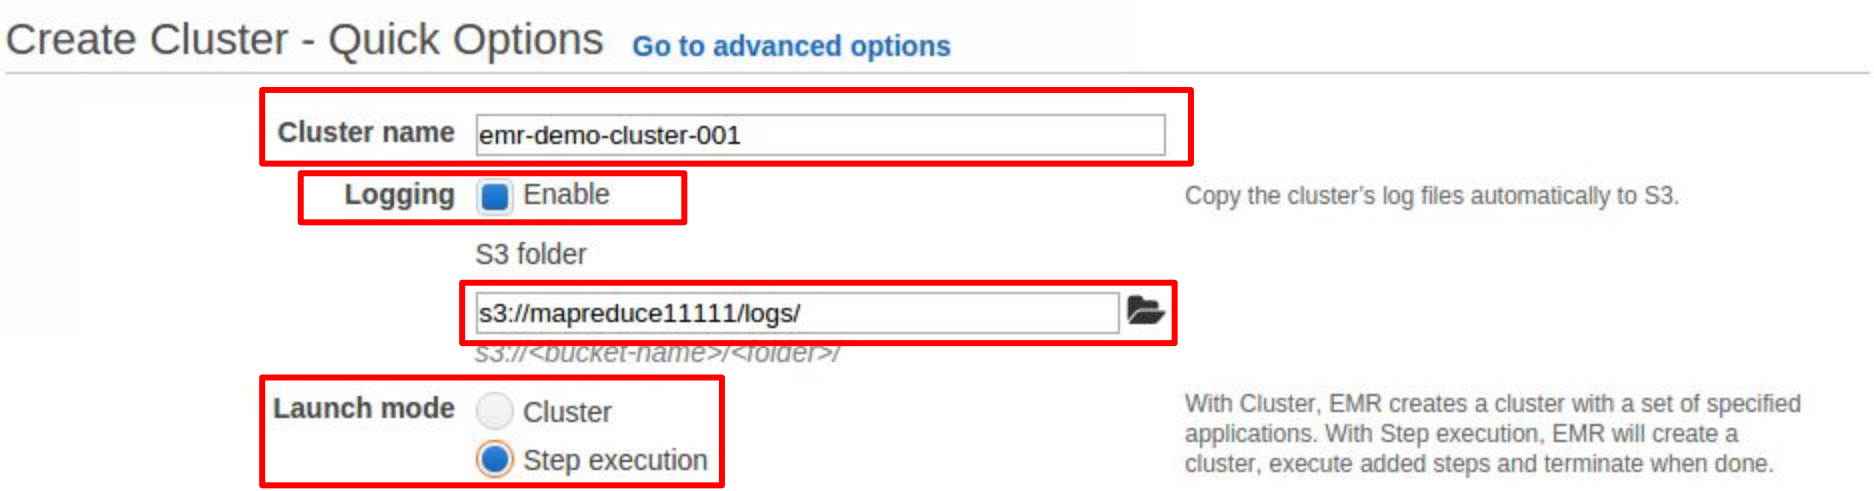

Streaming step: Ruby, Perl, Python, PHP, or Bash

Chose **Streaming program** from the **Step type** dropdown

Streaming step: Ruby, Perl, Python, PHP, or Bash

- ❏ Chose a unique name
- ❏ Select the mapper program from your **job** folder, the one that you just uploaded
- ❏ Select your reducer from the same location, or you can use the keyword **aggregate**. Amazon EMR supports the special aggregate keyword. For more information, go to the Aggregate library supplied by Hadoop.
- ❏ Next choose your input S3 location, which is the **input** folder from your bucket
- ❏ The output location is the **output** folder from your bucket followed by a **unique name** that you have to type in. In this case you can use **erm-python** after the output folder: **s3://MapReduce11111/output/emr-python**
- ❏ Press **Add**

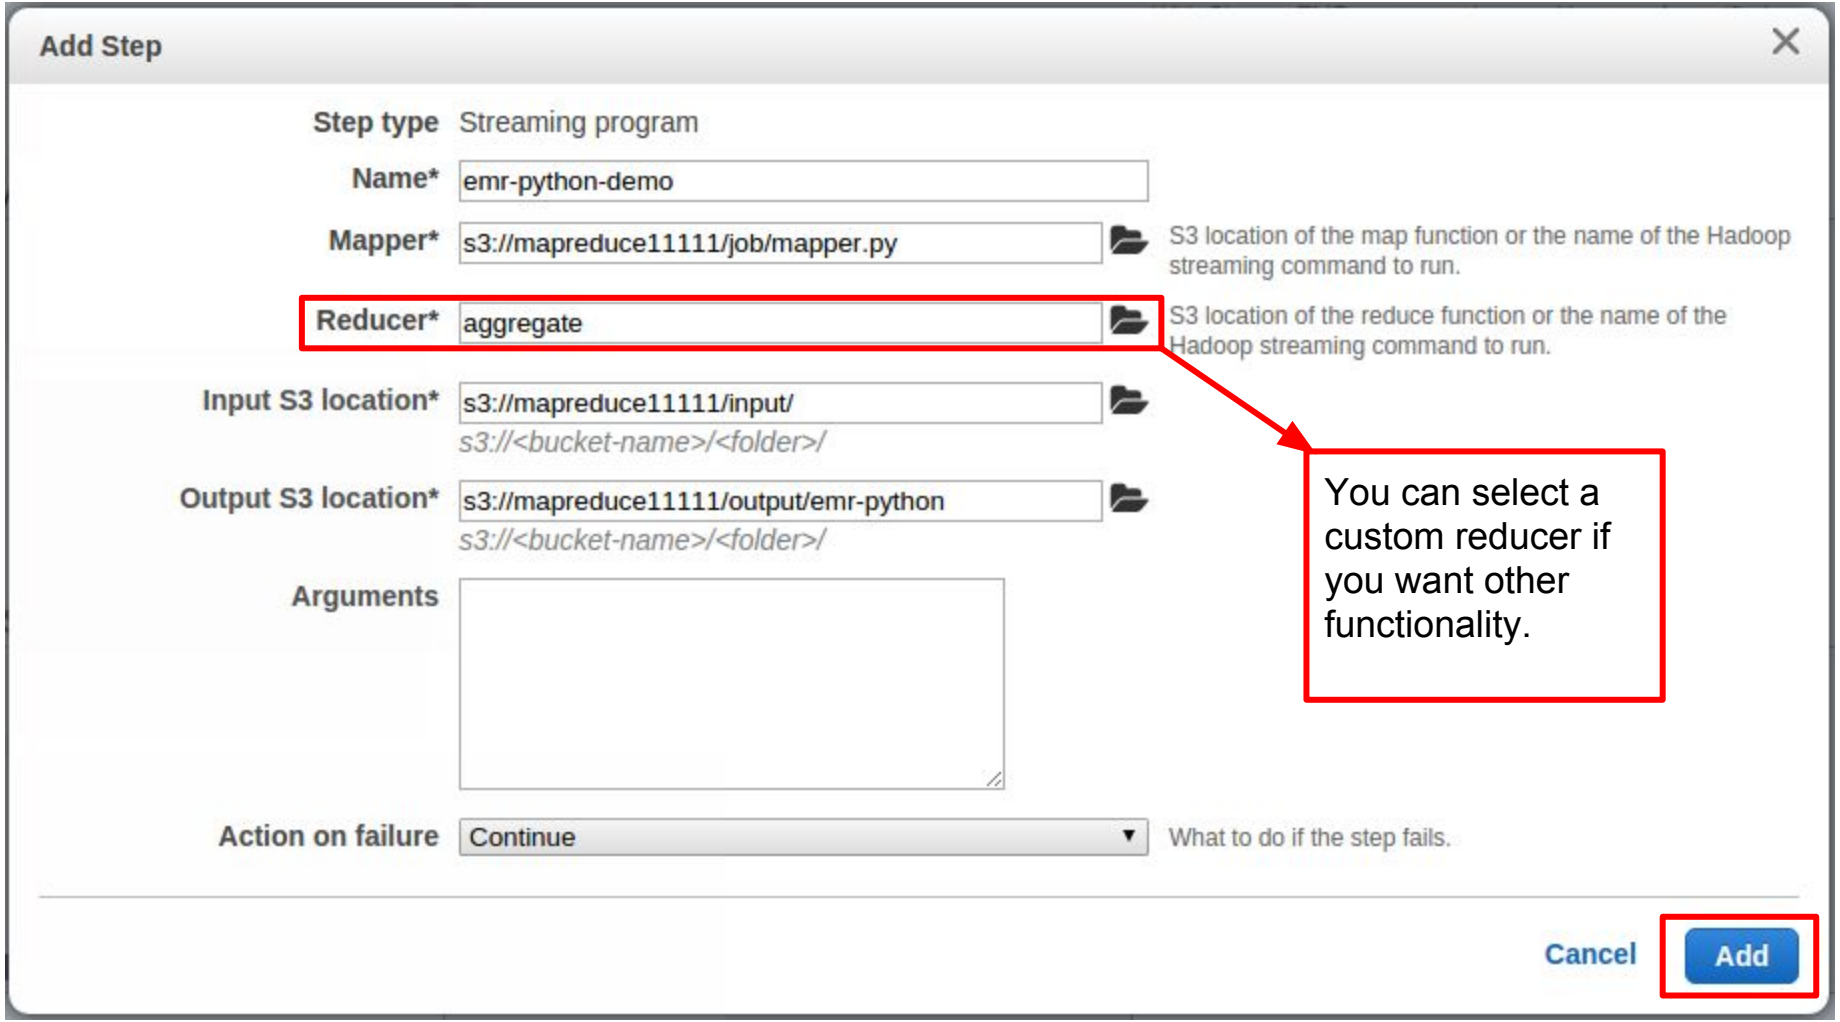
Add Step
Step type
Streaming program
Name*
emr-python-demo
Mapper*
s3://mapreduce11111/job/mapper.py
S3 location of the map function or the name of the Hadoop streaming command to run.
Reducer*
aggregate
S3 location of the reduce function or the name of the Hadoop streaming command to run.
Input S3 location*
s3://mapreduce11111/input/
s3://<bucket-name>/<folder>/
Output S3 location*
s3://mapreduce11111/output/emr-python
s3://<bucket-name>/<folder>/
Arguments
Action on failure
Continue
What to do if the step fails.
You can select a custom reducer if you want other functionality.
Cancel
Add

Software configuration

- Depending on how you chose to develop your MapReduce application, choose the appropriate **Vendor** and **Release**. For more info, please visit: http://docs.aws.amazon.com/ElasticMapReduce/latest/DeveloperGuide/emr-plan-hadoop-differences.html
- If you don't know what to choose, leave the default.

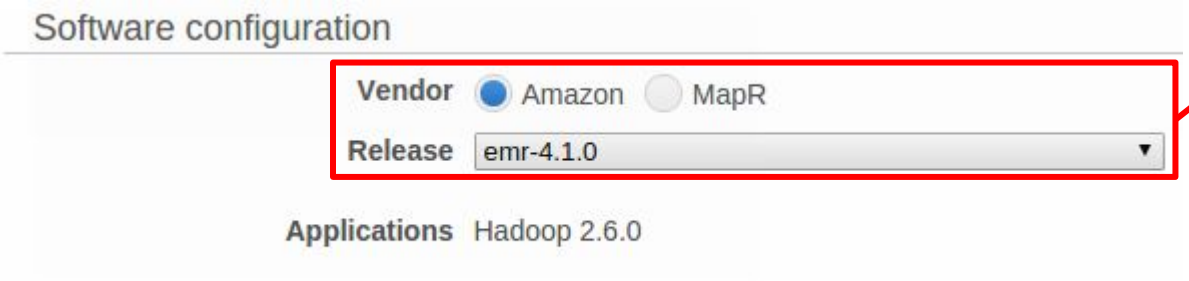

You can use the default (latest) version here, thus you do not need to change it with the version on this slide.

Hardware configuration and **Security**

- ❏ Choose your instance types, depending on your application requirements.
- ❏ Here is some information on the new instance types introduced in EMR: https://aws.amazon.com/blogs/aws/new-instance-types-for-amazon...

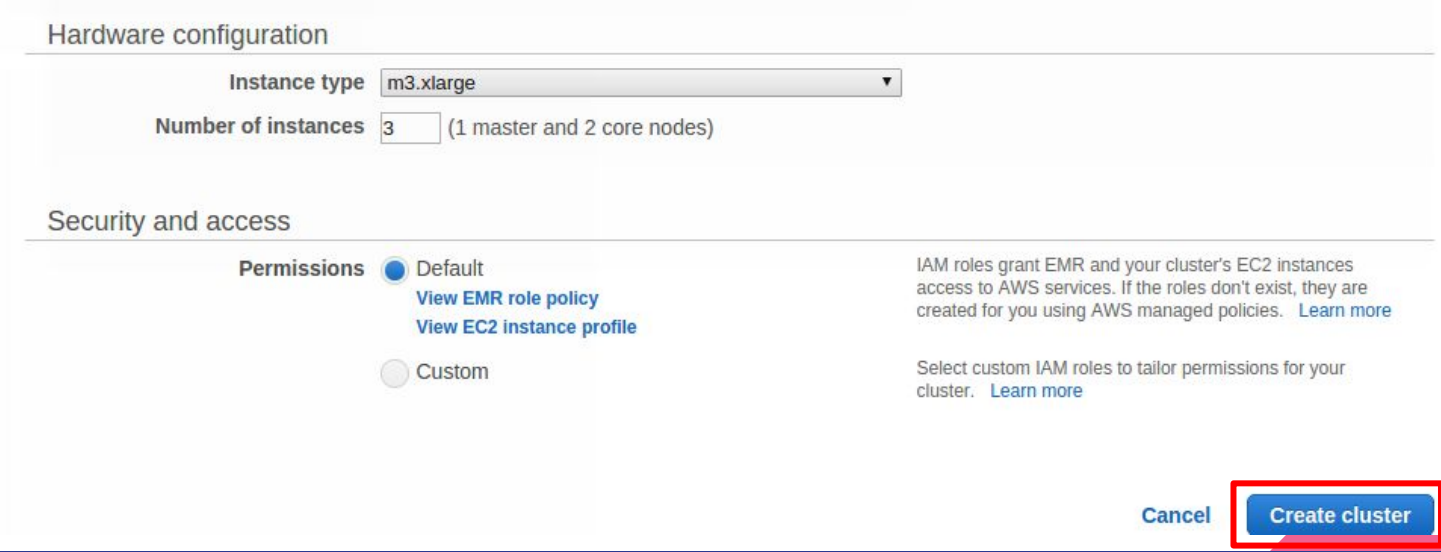

Check status and wait for completion

- Go to **Steps** under **Cluster details**, for the cluster that you just created
- Here you can see the progress of the cluster
- Refresh to update the status

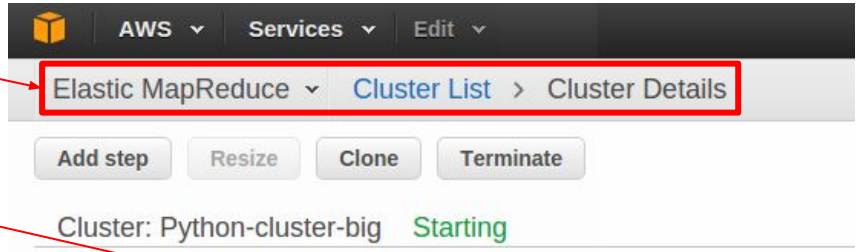

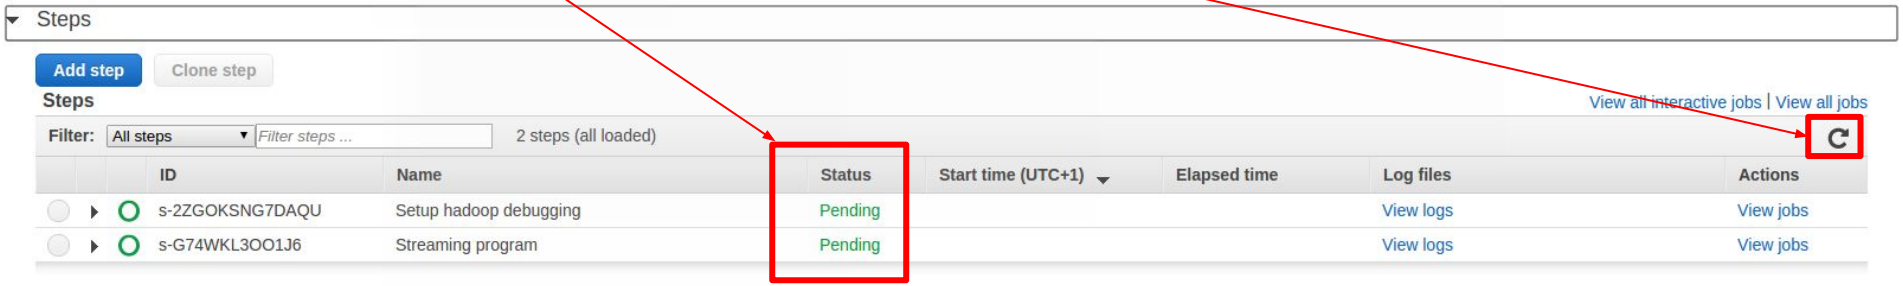

Wait until status changes to Completed

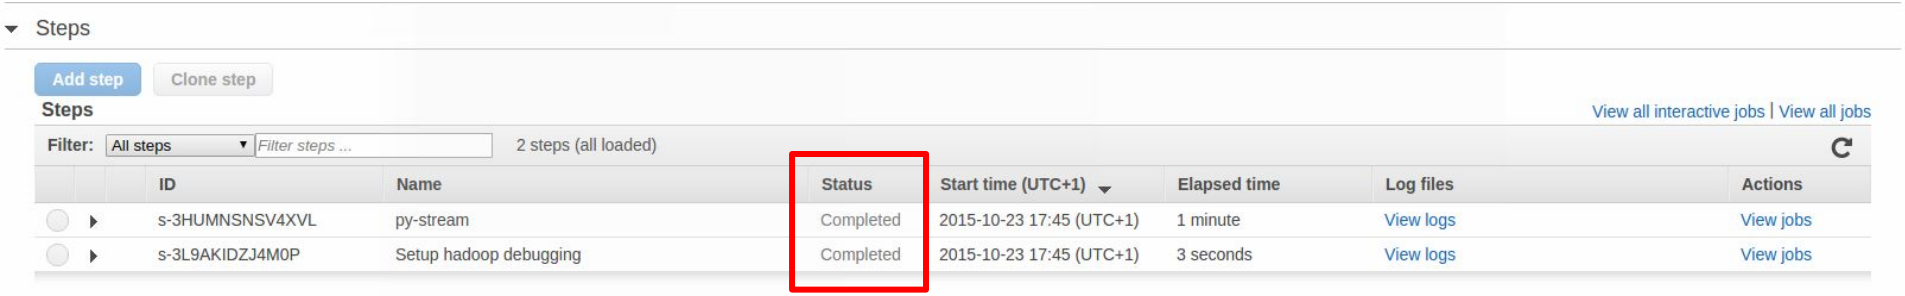

SUCCESS !!!

- Go back to the **Storage** and **Content Delivery > S3**
- Select your **S3 bucket**, the output folder and the unique name you chose when you created the cluster
- You should see here all output files

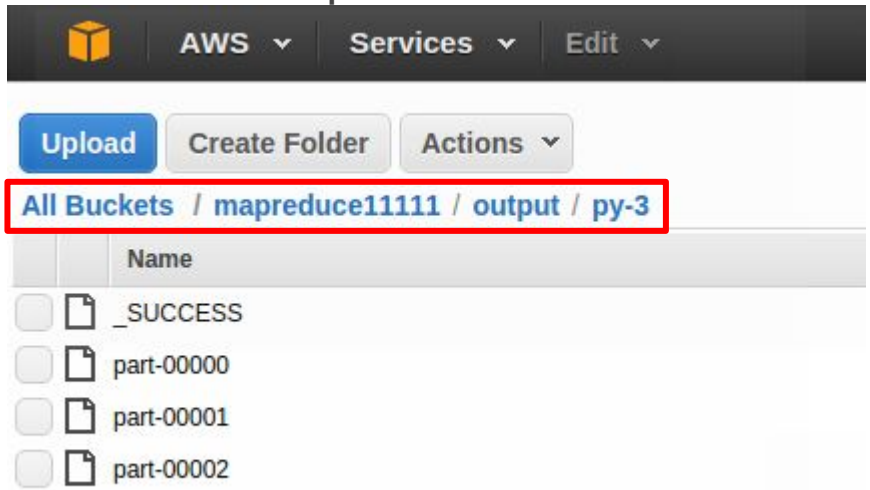

Download or process further

- You can now select the file that you want, press on Actions and choose Download
- Or you can reuse these in a new MapReduce program

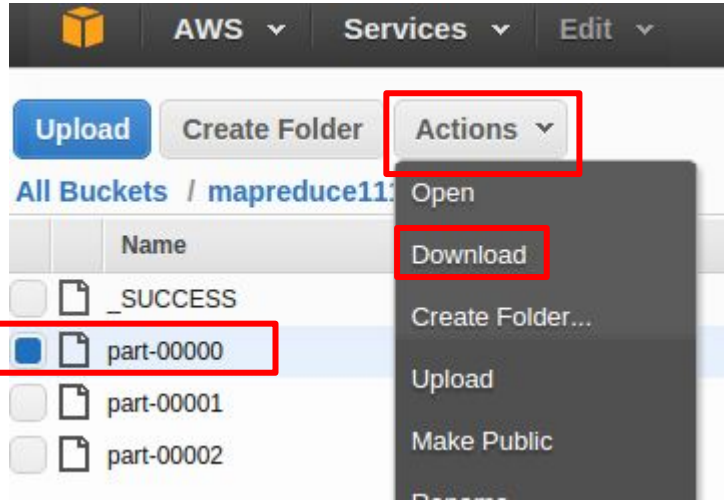

Congratulations, you have just successfully executed your first EMR program on AWS

Now please do the same for the google ngrams dataset, using the provided mapper and reducer.

For detailed EMR documentation, please visit: http://docs.aws.amazon.com/ElasticMap...

Java local development (your machine) in Eclipse

- For Java lovers, the following tutorial can help you get started: http://docs.aws.amazon.com/ElasticMapReduce/latest/DeveloperGuide/emr-common-programming-sample.html
  Please follow exactly all the steps to have the desired outcome.
- To carry out JUnit testing for your MapReduce code, please have a look at: https://cwiki.apache.org/confluence/display/MRUNIT/MRUnit+Tutorial
  which is a handy tip given by a fellow student Pavel Reich.

Useful links

- https://pythonhosted.org/mrjob/
- https://boto3.readthedocs.io/en/latest/
- http://hortonworks.com/products/sandbox/#downloads
- https://www.javacodegeeks.com/2015/03/running-pagerank-hadoop-job-on-aws-elastic-mapreduce.html